

Project 2: The Metal Part SORT-R

Foundations of Intelligent Systems, Fall 2015

Prof. Zanibbi

Due Date: Friday Nov. 20th, 2015 (11:59pm)

Submission Instructions


Through the dropbox in MyCourses, submit:

1. Your project write-up as a .pdf file, and
 2. A .zip archive containing:
 - Source code (.py) files (with main program files `trainMLP.py`, `executeMLP.py`, `trainDT.py` and `executeDT.py`.)
 - Data files needed to run the trained classifiers described in your report
 - A README file explaining how to run your programs and interpret their output
-

Problem Description

You own a start-up company that sorts and then sells used metal parts. As one of your first contracts, RIT has asked you to do this for parts taken from an old printing press. You need to sort the parts into **four categories: 1. bolts, 2. nuts, 3. rings, and 4. scrap**. After price negotiations with the school and some research of your own, you produce the table below defining the profit or cost in cents (\$0.01) obtained after sorting each part, based on the actual type of the part, and the class assigned automatically to the part by your system in development, SORT-R (patent pending).

Profit (in cents)				
Assigned	Actual			
	bolt	nut	ring	scrap
bolt	20	-7	-7	-7
nut	-7	15	-7	-7
ring	-7	-7	5	-7
scrap	-3	-3	-3	-3



Your system consists of a conveyer belt on which parts are carried, a camera and computer vision system that measure visual features for each part, and a sorter that places each part in a given bin. The vision system provides two measurements for each metal part:

1. **Six-fold rotational symmetry:** a measure of the average symmetry in the shape of the object, obtained by folding a black-and-white image in half in six different directions, each time counting the percentage of disagreeing pixels between the two sides of the fold. This helps to discriminate rings and nuts.
2. **Eccentricity:** a measure of variation in the distance to the center of the object as we trace the object's outer contour (roughly, a measure of how 'uncircular' an object is). This helps discriminate bolts from nuts and rings.

Your remaining task is to use machine learning algorithms to create a classifier that accurately identifies the type of each part as it arrives on the conveyor belt.

Data: CSV format training and testing datasets are provided. Both files contain one sample per line, with rotational symmetry in the first column, eccentricity in the second column, and the correct class (**given as 1 for bolt, 2 for nut, 3 ring and 4 scrap**) in the third column.

Multi-Layer Perceptron (MLP)

You will create programs to train and execute MLPs with one hidden layer.

Parameters: Network weights should be initialized randomly using values in the interval $[-1,1]$. Use the sigmoid (logistic) function as the node activation function, and a learning rate (α) of 0.1.

1. `trainMLP.py` takes a file containing training data as input and produces as output:
 - (a) **Five (5) files** containing the trained neural network weights after 0 (for initial weights), 10, 100, 1000, and 10,000 epochs.¹ Use batch training, repeatedly going over the training samples in-order, updating weights after each training sample is run. You will use these files to later run different network configurations on test samples using `executeMLP.py`, described below.
 - (b) An image containing a plot of the *learning curve*. The learning curve represents the total sum of squared error (SSE) over all training samples after each *epoch* (i.e. one complete pass over all training samples). Use the python matplotlib library (see <http://matplotlib.org/users/index.html>) to produce the plots.

On the CS computer systems, include `'from pylab import *'` at the top of your program to import matplotlib.

2. `executeMLP.py` takes a file defining a trained network and a data file, and runs the network on the data file. The program should produce:
 - (a) **(standard output)** On the command line:
 - i. Recognition rate (% correct)
 - ii. Profit obtained
 - iii. A *confusion matrix*, which is a histogram counting the number of occurrences for each combination of assigned class (rows) and actual class (columns). The main diagonal contains counts for correctly assigned samples, all remaining cells correspond to counts for different types of error.
 - (b) **(image file)** A plot of the test samples (using shapes/color to represent classes) along with the *classification regions*. The program runs the current network over a grid of equally spaced points in the region (0,0) to (1,1) (the limits of the feature space), using a different color to represent the assigned class (i.e. classification decision by the MLP) at each point. Use matplotlib to produce this image.

Decision Tree (DT)

We will use a modified form of decision tree that allows *continuous* features to be partitioned using binary splits (similar to what is used in the C4.5 algorithm). To find the best split for a

1. You may use the Python 'pickle' module, to easily save and recover object states, or use a simple text file (e.g. CSV file) representing the MLP network architecture and weights.

continuous feature, we: 1) sort samples in increasing order by their values for one attribute, and then 2) compute the information gain at every possible split point between adjacent samples, using their midpoint as the split value s . Samples with an attribute value $\geq s$ are then put in the ‘right’ child of the current node, and otherwise the left child of the current node. **Note that unlike with discrete attributes, we will need to re-use features.**

We will also apply Chi-Squared pruning to the learned tree, to avoid overfitting. For the Chi-Squared test, try using a significance level of 1% or 5%.

You need to produce two programs for your decision trees.

1. `trainDT.py` takes a training data file and produces:
 - (a) **Two files** representing the trained decision tree 1) after automatic induction, and 2) after subsequent pruning using the Chi-Squared test. You will later run these trees using `executeDT.py`, described below.²
 - (b) **Two images illustrating splits in the two decision trees.** For each tree, use matplotlib to create a plot visualizing how the feature space is partitioned by the decision tree, by drawing a box around both regions created at every internal (‘split’) node in the tree. **You should also show the training samples in the plots.**
 - (c) On the terminal, print a summary of the contents for each decision tree:
 - i. Number of nodes, and number of leaf (decision) nodes
 - ii. Max, minimum and average depth of root-to-leaf paths in the decision tree
2. `executeDT.py` takes a file defining a decision tree ensemble and a data set, and produces the same output as `executeMLP` for decision tree in the file.

Experiment

Use `trainNN.py` and the `train_data.csv` data file to train a multi-layer perceptron with two input nodes, five hidden nodes, and four output nodes (for four classes), plus bias nodes. Then, use `executeNN.py` to run each network (all five versions) on `test_data.csv`. Collect the performance metrics and classification region images for each classifier. Repeat this a number of times (**remembering to save your previous results!**), and then report the results for what you consider to be the best run.

Then run `trainDT.py` to create your decision trees from `train_data.csv`, and then use `executeDT.py` to run both trees over `test_data.csv`.

Write-Up

1. **Algorithms:** Briefly discuss the similarities and differences between the two learning algorithms. Which type did you expect to perform better in the experiment, and why?
2. **Data:** Provide separate plots for the training and test data sets. Show sample classes using colors and/or shapes. Comment on the distribution of classes in the data sets.
3. **Results:** Provide the following for the test data results.
 - (a) MLP:
 - i. Plots showing the test samples and classification regions produced by different numbers of training epochs

2. For the decision trees, pickling is probably simplest.

- ii. The learning curve image (SSD vs. Epoch) for the trained MLP
 - iii. A table showing the recognition rate and profit for each number of saved epochs for the MLP
 - (b) Decision Trees:
 - i. Plots showing the test samples and classification regions produced by each of the two decision trees.
 - ii. Plots showing how feature space is split by each decision tree.
 - iii. A table providing the recognition rate and profit obtained by each decision tree, along with the tree metrics produced by `trainDT.py`.
4. **Discussion:** Comment on the results (in prose, not bullet points!), making sure to answer the following questions:
- Which versions of the classifiers performed best in terms of 1) accuracy and 2) profit? Did this meet your expectations?
 - How do the hypotheses (i.e. class boundaries) and performance metrics differ between the different version of the MLP and decision trees, and why?

Use your result figures and tables, along with the confusion matrices to provide evidence for your arguments. Also provide any additional comments on the problem, classifiers, results or process that you feel is informative or interesting.

Grading

- 40% Correctness of implementation
- 10% Test set accuracy for the best classifiers
- 10% Coding style (use a research programming style)
- 40% Write-up
- (5%) **Bonus:** try using different numbers of hidden nodes in the MLP and/or different Chi-Squared pruning thresholds, and include any interesting results and discuss them in your report.

Acknowledgement

This dataset and part image is taken from the introductory pattern recognition textbook, "Classification, Parameter Estimation and State Estimation" by van der Heijden, Duin, Ridder and Tax (Wiley, 2004).